WikiPortraits card report

Generated 2026-10-13 17:10 UTC · model llm-clef-flash on Wikimedia LiftWing · 1 of 1 images scored

Method

Model: llm-clef-flash (Cloudflare/clef-flash, 9.41 B params BF16, Apache-2.0) served by Wikimedia LiftWing at api.wikimedia.org/service/lw/inference/v1/models/llm-clef-flash:predict (KServe on HuggingFace transformers, not vLLM).

Schema: one state + six typed questions, all answered in a single forward pass; the model generates no text (usage.output_tokens = 0). Question types: face_visible (yes/no), pose (choice of 5), quality (score 0-4), infobox_ready (yes/no), publish_ready (choice of 2: publish / pass), eyes_to_camera (yes/no).

Headline numbers: composite = min(face_visible, eyes_to_camera) x pose weight (head-and-shoulders 1.0, three-quarter 0.9, full body 0.6, candid 0.5, group 0.25). That is our arithmetic on three model answers, NOT a model probability; P(publish) and P(infobox) are the model's own. Ranking uses the composite.

Spread: the run reports candidate spread (range and median gap among answers at or above 0.25), not standard deviation — sd rises whenever answers pile up at either end, including at the ceiling, so it rewards the wrong thing.

Inputs: Commons thumbnails at the standard step nearest above 640 px (the server downscales to 640 px wide before the encoder, so larger sends add no information).

Note: probabilities depend on the question set — the head scores every question jointly, so only compare runs that used this identical five-question schema.

Summary

Images pasted 1 · resolved 1 · scored 1 · not scored 0

P(infobox-ready): mean 0.6496 · median 0.6496 · max 0.6496 · 1 of 1 at or above 0.5

P(publish): mean 0.9155 · composite (ours, min(face, eyes) x pose weight): mean 0.9406, max 0.9406

Quality (0-4): mean 2.56

Pose mix: head_shoulders 1

Wall: median 56.735 s per call, 56.73 s total · input tokens 1323 · job wall clock 0.13 s

CANDIDATE SPREAD (composite >= 0.25): 1 of 1 photos are candidates · range None · median gap None · 0 of 1 (0%) in the uninformative 0.35-0.65 band

same measure for infobox_ready: 1 of 1 candidates, range None, median gap None, 1 in the band

same measure for publish_ready: 1 of 1 candidates, range None, median gap None, 0 in the band

Endpoint latency (service-side): median 56.735 s per call (56.735-56.735 s, n=1) vs a measured warm baseline of 0.54 s - 105.1x

All images, ranked by composite

#	Photo	composite	P(publish)	P(infobox)	pose	quality	P(eyes)	wall
1	MKr377445 Susanne Wuest (In die…	0.941	0.915	0.650	head_shoul…	2.56	0.941	56.73s

Per-image determinations, with the photo the model saw



MKr377445 Susanne Wuest (In die Sonne schauen, Cannes 2025).jpg

composite 0.941 (ours) · P(publish) 0.915 (publish) · P(infobox) 0.650

face_visible 0.971 · eyes_to_camera 0.941 · pose head_shoulders

pose head_shoulders 0.946 three_quarter 0.025 candid 0.021 full_body 0.005 group 0.003

quality 2.560 of 4 · 0 0.040 1 0.063 2 0.325 3 0.440 4 0.132

original 3465x5197 px, sent 960x1440 · 1323 input tokens · 56.735 s · HTTP 200

Licence: CC BY-SA 4.0 · photo: Martin Kraft

commons.wikimedia.org/wiki/File:MKr377445_Susanne_Wuest_(In_die_Sonne_schauen,_Cannes_2025).jpg

Every probability above was returned by the endpoint; nothing is inferred. Photographs are from Wikimedia Commons under their own licences. Verify before any on-wiki use: model output is an assistive draft, not a source.